

Machine Learning for Large-Scale Data Analysis and Decision Making (MATH80629A) Winter 2023

Week #12- Summary

Announcement

- **Last Quiz, Quiz#6 on RL: Next week**
- **Project Presentation: in-person on April 14**
- **Project Report: due April 23**
- **Final exam: April 30**



Today

- Summary of Sequential decision making I
- Q&A
- Finish RSs case study
- Hands on session

Sequential decision making I

Reinforcement Learning

- Sometimes we need a model where **the learning** and **the decision making** interact closely
- Imagine building a robot that must navigate autonomously
 - The robot has wheels and a camera
- You think about using a **two-stage approach**:
 1. Use supervised learning to identify objects in scenes
 2. Given scene content have a decision-making module that controls its wheels


Limitations of two-stage approach

- **Supervised learning doesn't know about the decision-making**
 - Its objective is, for example, to maximize accuracy
- **For decision making, different errors have different costs**
 - E.g., missing the cliff could have dire consequences. missing sky less so.
 - Incorporating these costs into the learning objective is tough
- **Several other limitations:**
 - need labeled data
 - improvements in SL do not necessarily lead to improvements in decision making
 - ...

Alternative: Reinforcement learning (RL)

- **Incorporates both stages in a single framework**
- Incorporates the ideas of:
 - **state (observation)**
 - **action**
 - **reward**

Initial example with grid world



- Each cell is a state (S)
- Actions indicate which movements are possible:

$$A := \{L, R, U, D\}$$


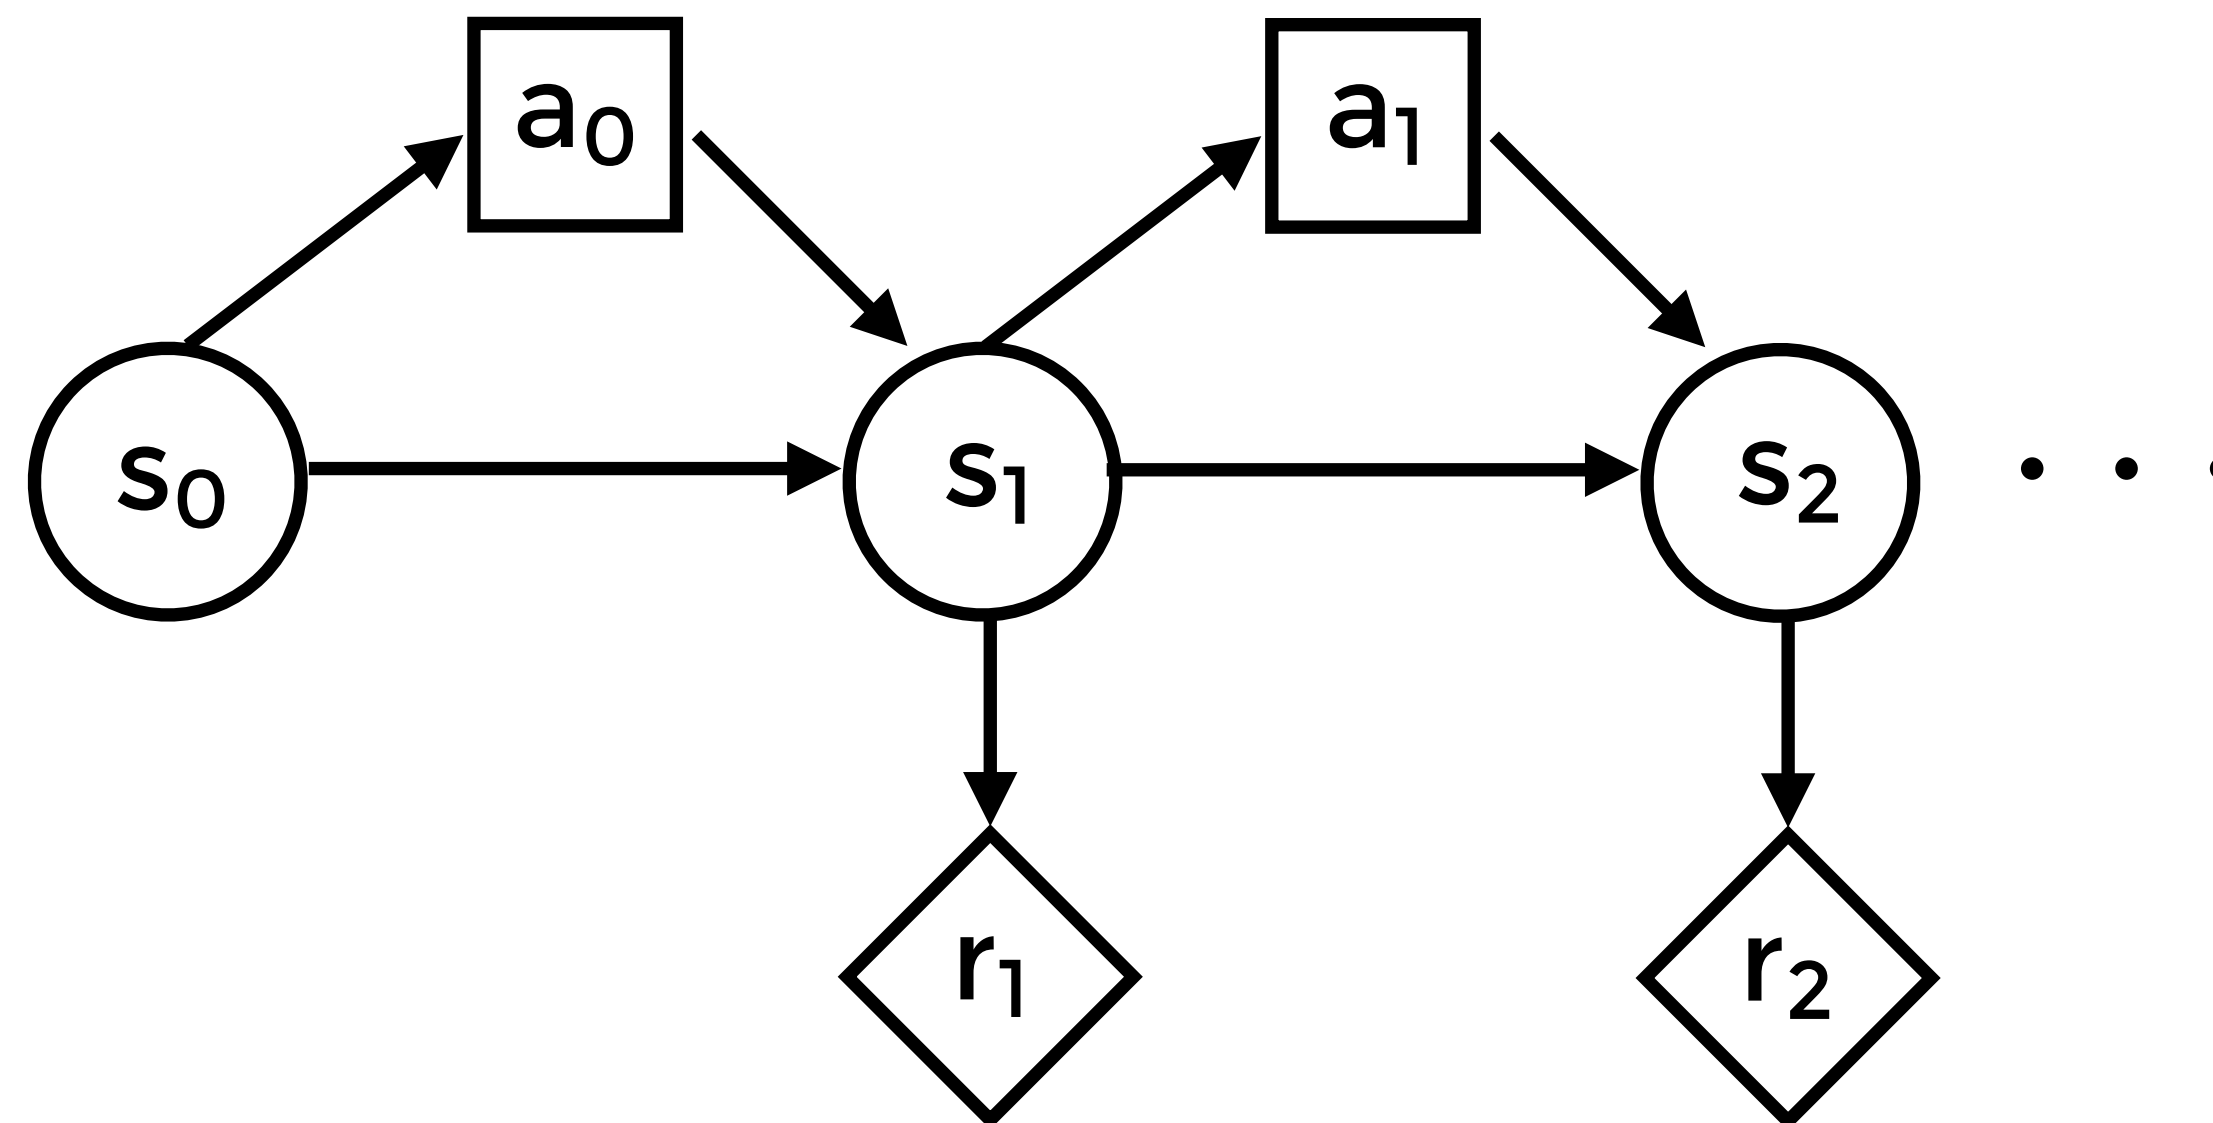
- Rewards encode the task:
 $R(s)$
- Transition probabilities encode the outcome of an action:
 $P(s' \mid s, a)$

This week we discuss a version of RL where these are observed

Markov Decision Process (MDP)

- Provide a framework for decision-making under uncertainty
- Markov process with decisions and utilities
- Assumes stationarity (i.e., transitions are fixed across time)


- Square nodes: decisions
- Circle nodes: States
- Diamond nodes: utility



Markov Decision Process (MDP)

$$\langle A, S, P, R, \gamma \rangle$$

- **A**: set of actions
- **$P(S' | S, A)$** : transition probabilities
- **$R(S)$** : reward function
- **γ : discount factor** $\in [0, 1]$
- **A policy**: $\pi : S \rightarrow A$
- **Goal**: find the optimal policy



Optimal policy?

- **Agent is trying to maximize its rewards (utility)**
- **Utility simply assigns a real value to a state**
- **Typically combine rewards with an additive function**

$$\sum_t R(s_t)$$

Discounting (γ)

- The sum of rewards could be infinite/unbounded

$$\lim_{T \rightarrow \infty} \sum_t^T R(s_t)$$

- A typical solution is to use a discount factor $0 \leq \gamma \leq 1$

$$\lim_{T \rightarrow \infty} \sum_t^T \gamma^t R(s_t)$$

- Geometric series. Bounded by: $\frac{R_{\max}}{1 - \gamma}$
- Intuition: would rather have rewards sooner

Solving an MDP

- Find the optimal policy of an MDP

$$\pi^*(s) \quad \forall s$$

- Policies are evaluated using their expected utility:

$$EU(\pi) = \sum_{t=0}^{\infty} \gamma^t \sum_{s_{t+1}} P(s_{t+1} \mid s_t, \pi(s_t)) R(s_{t+1})$$

- The optimal policy is the one with highest expected utility:

$$EU(\pi^*) \geq EU(\pi) \quad \forall \pi$$

Solving an MDP

- 1. Value iteration**
- 2. Policy Iteration**

Value Function

- $V(s_t)$: The value of being in state s at time t

$V(s_t) :=$ expected sum of rewards of being in s

Finite horizon

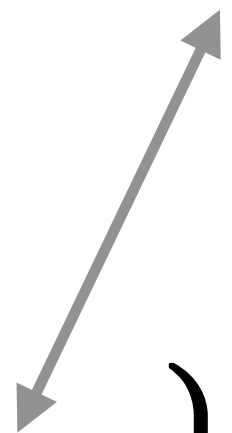
- Assume that the process has T steps

- The value at step T is $V(s_T) = R(s_T)$

- The value at step $T-1$ is

$$V(s_{T-1}) = \max_{a_{T-1}} \left\{ R(s_{T-1}) + \gamma \sum_{s_T} P(s_T | s_{T-1}, a_{T-1}) R(s_T) \right\}$$

- The value at step t is $(0 \leq t \leq T)$

$$V(s_t) = \max_{a_t} \left\{ R(s_t) + \gamma \sum_{s_{t+1}} P(s_{t+1} | s_t, a_t) V(s_{t+1}) \right\}$$


Bellman equation

- Value of state s

$$V(s_t) = \max_{a_t} \left\{ R(s_t) + \gamma \sum_{s_{t+1}} P(s_{t+1} \mid s_t, a_t) V(s_{t+1}) \right\} \quad \forall s$$

- Recursive equations
- The value of a state only depends on the state's reward and the neighbours' value
- This is also known as a dynamic programming equation

Value iteration (VI)

- Iteratively update $V(s)$ for each state until convergence
- (Initialize $V(s)$ for every state)

- For $i=1,2,3,\dots$

- For $s=1,\dots,S$

$$V(s) = \max_a \left\{ R(s) + \gamma \sum_{s'} P(s' | s, a) V(s') \right\}$$

- The policy is implicit

- Once converged: $\pi^*(s) = \arg \max_a \left\{ R(s) + \gamma \sum_{s'} P(s' | s, a) V^*(s') \right\} \quad \forall s$

Policy Iteration (PI)

- Improve policy explicitly.

Start with any (e.g., random) policy π

Iterate until convergence:

1. Given current policy get the value of each state

$$\mathbf{V}^{\pi}(\mathbf{s}) = \mathbf{R}(\mathbf{s}) + \gamma \sum_{\mathbf{s}'} \mathbf{P}(\mathbf{s}' | \mathbf{s}, \pi(\mathbf{s})) \mathbf{V}^{\pi}(\mathbf{s}') \quad \forall \mathbf{s}$$

2. Update the current policy

$$\pi'(\mathbf{s}) = \arg \max_{\mathbf{a}} \left\{ \mathbf{R}(\mathbf{s}) + \gamma \sum_{\mathbf{s}'} \mathbf{P}(\mathbf{s}' | \mathbf{s}, \mathbf{a}) \mathbf{V}^{\pi}(\mathbf{s}') \right\} \quad \forall \mathbf{s}$$

Policy
Evaluation

Policy
Update

PI vs. VI

- **Value iteration is faster per iteration**
- **Policy iteration converges in fewer iterations**

MDP Real-world Examples

Examples are taken from:

<https://towardsdatascience.com/real-world-applications-of-markov-decision-process-mdp-a39685546026>

MDP framework

- To express a problem using MDP, we need to define the followings
- **states** of the environment
- **actions** the agent can take on each state
- **rewards** obtained after taking an action on a given state
- **state transition** probabilities.

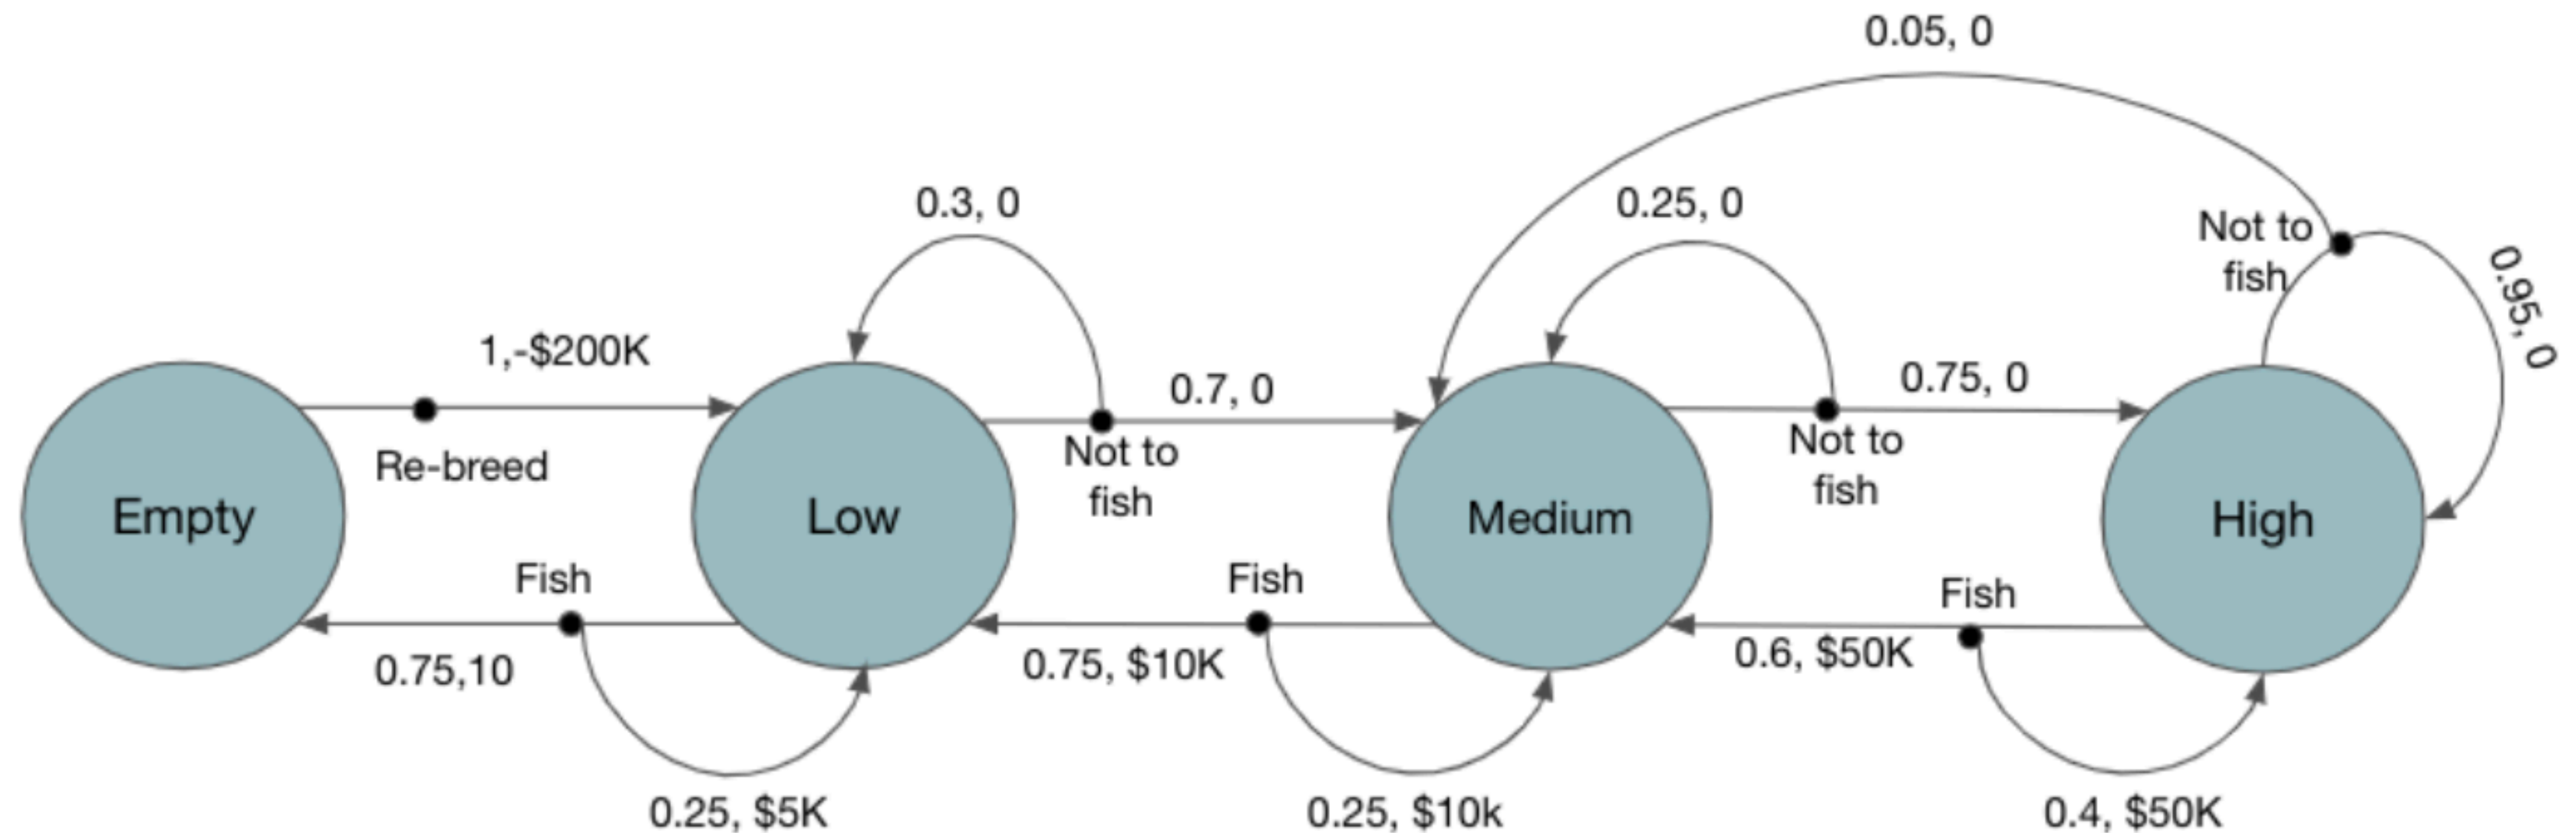
Salmon Fishing



- **States:** The number of salmons available in that area in that year. E.g., four states; **empty**, **low**, **medium**, **high**.
- **Actions:** **fish** and **not_to_fish**. Fish means catching certain proportions of salmon. For the state empty the only possible action is not_to_fish.
- **Rewards:** Fishing at certain state generates rewards, let's assume the rewards of fishing at state low, medium and high are \$5K, \$50K and \$100k respectively. If an action takes to empty state then the reward is very low -\$200K as it require re-breeding new salmons which takes time and money.

Salmon Fishing

State Transitions: Fishing in a state has higher a probability to move to a state with lower number of salmons. Similarly, not_to_fish action has higher probability to move to a state with higher number of salmons (excepts for the state high).

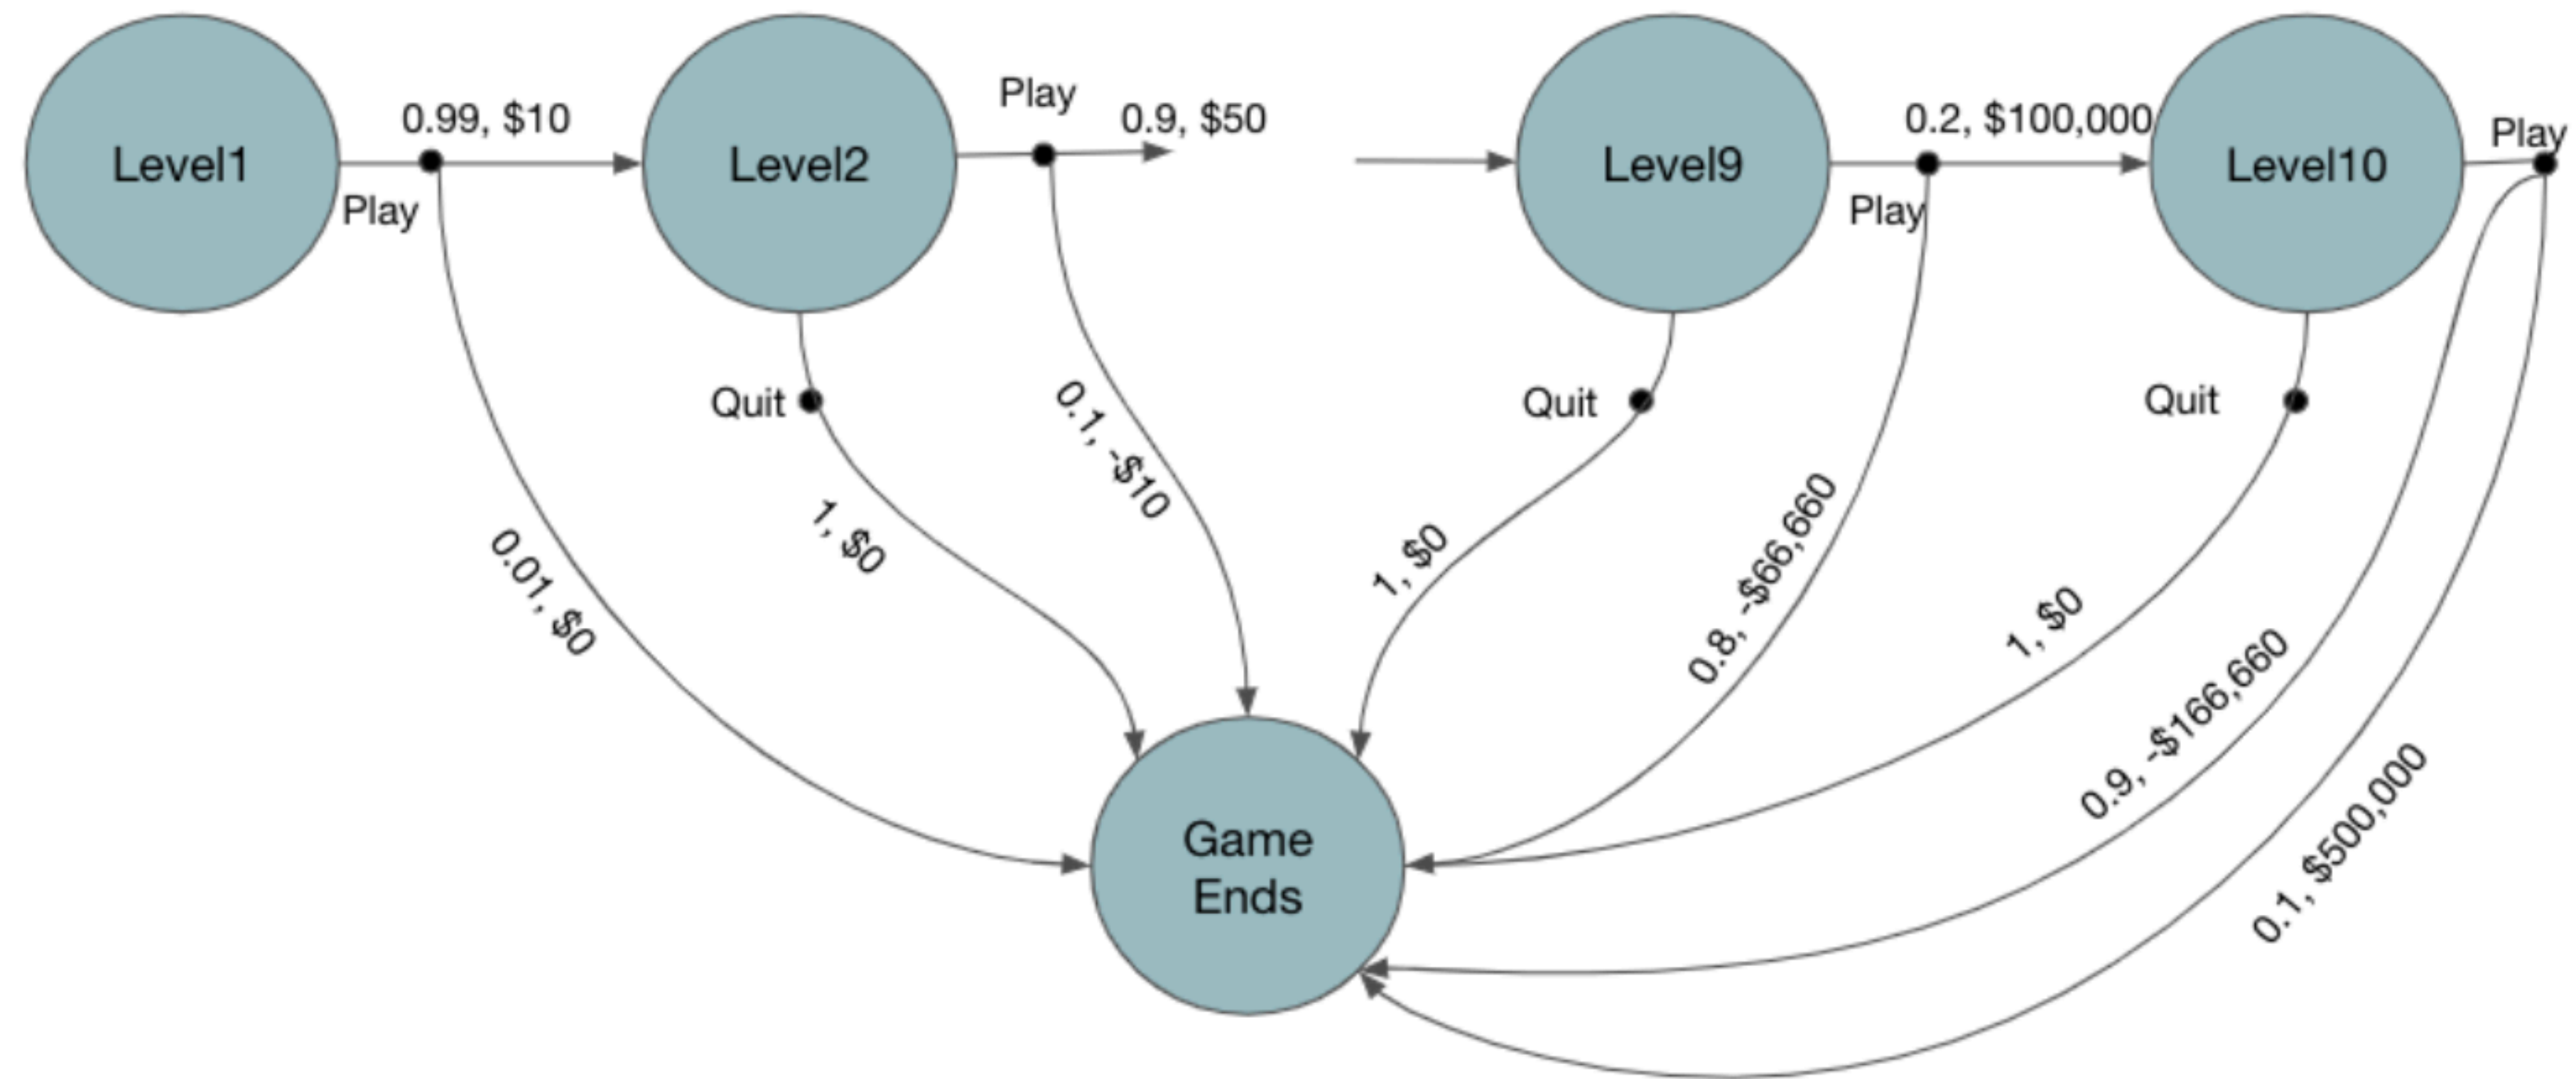


Quiz Game Show



- **States:** {level1, level2, ..., level10}
- **Actions:** Play at next level or quit
- **Rewards:** Play at level1, level2, ..., level10 generates rewards \$10, \$50, \$100, \$500, \$1000, \$5000, \$10000, \$50000, \$100000, \$500000 with probability $p = 0.99, 0.9, 0.8, \dots, 0.2, 0.1$ respectively. The probability here is a the probability of giving correct answer in that level. At any level, the participant losses with probability $(1 - p)$ and losses all the rewards earned so far.

Quiz Game Show



Other MDP examples

- **Harvesting**: how much members of a population have to be left for breeding.
- **Agriculture**: how much to plant based on weather and soil state.
- **Water resources**: keep the correct water level at reservoirs.
- **Inspection, maintenance and repair**: when to replace/inspect based on age, condition, etc.
- **Purchase and production**: how much to produce based on demand.
- **Queues**: reduce waiting time.
- **Finance**: deciding how much to invest in stock.
- **Robotics**: navigator, dialogue system, etc.

Other MDP examples

- Interested to find more? Check White, D.J. (1993) that mentions a large list of applications.